

August 7, 2019

Dear reviewers,

Thank you very much, again, for your comments on our manuscript “Evaluative adjective sentences: A question-based analysis of projection”. We have now resubmitted our manuscript to Glossa. This new manuscript has taken into consideration your comments, as we detail in this letter (your comments in black, our responses in blue).

We have made the following larger changes to the manuscript in response to the reviewers’ comments:

1. In response to comments by reviewer A, section 1 now better contextualizes the analysis developed in section 2. Specifically, section 1 now already introduces the reader not just to analyses that derive projection from lexical specification of the content as a presupposition but also to analyses that derive from projection from the discourse status of the content. This section also introduces the reader to motivation for developing such analyses, contents for which such analyses have been proposed, and motivation behind deriving the projectivity of the prejacent of EAS from its discourse status. This section also clarifies what it means for content to be less projective than other content, and that content being lexically specified as a presupposition is still a viable option for other projective content.
2. In response to comments by reviewers A and B, we have expanded section 2.1 to give the reader more background on Simons et al’s 2017 question-based analysis of projection and on Tonhauser et al’s 2018 modification, according to which content projects to the extent that it is not at-issue.
3. In response to review B, the second paragraph of section 2.4 now explicitly compares our projection analysis to the ambiguity analysis presented in Karttunen et al 2014.
4. The first paragraph of section 5 now briefly discusses implications for theories of projection, in response to comments by reviewers A and C.

Kind regards,

The authors

Review A:

Summary: This paper adduces experimental data to argue against the view that the prejacent of EAS is a lexically specified presupposition, instead claiming that the prejacent is a lexical entailment that projects just when it is not at-issue.

Overall assessment: I believe this paper has the potential to make a valuable contribution to the field. At the same time, I find (i) there are often swiping claims that obscure what has actually been achieved and (ii) there is a lot of confusion about theory, data, and mostly how the two connect. I elaborate on these points below.

#1: The authors claim to show that the prejacent of EAS is not a lexical presupposition but rather a lexical entailment of EAS. (The crucial term “lexical entailment” is never defined in the paper, so I take it to simply mean an entailment of an expression that lacks any projection triggers.)

Thank you for this remark: we now clarify in footnote 7 of section 1 that we follow Abrusán 2011 in our use of the term ‘lexical entailment’: lexical entailments are entailments of sentences, not predicates, and are distinguished from logical entailments. So, yes, the prejacent and the generalization are lexical entailments and we argue that their projectivity is better derived from their discourse status than from lexically specifying these contents as presuppositions.

Footnote 9 now clarifies that evaluative adjectives also trigger subject-related presuppositions (those are discussed in Barker 2002). Thus, it is not the case that a lexical entailment is an entailment of an expression that does not trigger presuppositions.

In reality however, all that is shown is that the prejacent of EAS is not highly projective (although somewhat different from presuppositions associated with “soft” triggers like know) and that it projects just when not at-issue.

Right, we observe that the prejacent of EASs is not highly projective. As we now clarify in section 1, we take this observation, as well as the observed interaction in the projection of the prejacent and the generalization, as impetus to developing an analysis according to which the prejacent is not a lexically specified presupposition.

But discourse status alone doesn’t explain why the prejacent of EAS is able to project in the first place, given that lexical entailments never do that.

As mentioned at the beginning of this letter, we now provide more background on analyses that derive projection from the discourse status of utterance content in sections 1 and 2.1.

The authors cite the Projection Principle of Tonhauser et al. 2018:499 below, but the formulation overstates Tonhauser et al.’s actual claim, which only about projective contents, or contents that in general have the potential to project.

(1) Gradient Projection Principle: If content C is expressed by a constituent embedded under an entailment-canceling operator, then C projects to the extent that it is not at-issue.

We disagree with the claim that our formulation of the Gradient Projection Principle overstates Tonhauser et al.’s actual claim: the Gradient Projection Principle as given in our manuscript is taken verbatim from Tonhauser et al 2018:499, as shown by the screen shot below. Furthermore, contrary to what is claimed in this part of the review, the Gradient Projection Principle, as stated in Tonhauser et al 2018, applies to any content C

that is expressed by a constituent embedded under an entailment-canceling operator. It is not restricted to lexical entailments but applies to ANY content.

judgment that a content is projective to a certain extent reflects the probability with which they believe the speaker to be committed to the content. On this interpretation, speaker commitment may be a binary, categorical property and projection variability arises from the listener's uncertainty about the whether the speaker is committed. In this paper, we remain agnostic about the underlying interpretation of projectivity as a gradient property, though our discussion of projection variability will be in line with the first interpretation.

Given the hypothesis that projective content varies in how projective it is, we propose the Gradient Projection Principle in (7):

- (7) **Gradient Projection Principle:** If content C is expressed by a constituent embedded under an entailment-canceling operator, then C projects to the extent that it is not at-issue.

The point can be made quite clearly on the following example (a version of what is cited in the paper on pp.6-7). In B's response, the complement does not answer the Discourse Question (only the matrix does), so the complement is not at-issue, but it projects only with discover, not with think. If (1) were allowed to apply indiscriminately to all contents (including lexical entailments), the contrast below would be a puzzle, so the above principle needs to be circumscribed to projective content.

(2) Context: Henry and Harriet are an academic couple that lives on the West Coast.

A: Why is Henry in such a bad mood?

B: He discovered/thinks that Harriet was at Princeton for a job interview.

Contrary to what is assumed here (and also implicitly in many parts of the literature), it is not the case that non-entailed content like the content of the complement of non-factive predicates like *think* does not project. First, there are discussions in the published non-experimental literature about the projection of non-entailed content, like the content of the complement of non-factive predicates: e.g., Schlenker 2010 (cited), Anand & Hacquard 2014 (The Art and Craft of Semantics: A Festschrift for Irene Heim), Spector & Egré 2015 (A uniform semantics for embedding interrogatives: An answer, not necessarily the answer; *Synthese*), Simons et al 2017 (cited).

Second, there is experimental research that shows that non-entailed content can project: e.g., de Marneffe et al 2019 (*Sinn und Bedeutung* proceedings) show that the content of the complement of many non-factive predicates projects in naturally occurring data, some even as strongly as that of some factive predicates. Tonhauser & Degen (2019, talk presented at Experimental Pragmatics in Edinburgh) likewise show experimentally that non-entailed content is projective. Finally, Tonhauser et al 2018 included non-entailed content, like the content of the complement of *confess* and *establish*, and show not only that it is projective but also that its projectivity is predicted by its at-issueness.

In sum: We know from previous literature that the prejacent of EAS can project, and Karttunen et al. 2014 tell us that projection need not arise. The novel contribution of the current paper is the empirical observation that the projection of an EAS prejacent is correlated with the lack of at-issueness. Still, we are offered no mechanism that explains what makes the prejacent of EAS able to project in the first place, so we are left in the dark on the nature of this inference. Whether or not presuppositional, it is clearly not a lexical entailment.

Thank you, again, for asking us to clarify the use of the term ‘lexical entailment’. Regarding the claim that “we are offered no mechanism that explains what makes the prejacent of EAS able to project in the first place”, we have now, as mentioned above, included more background on analyses that derive projection from discourse status.

#2: The paper claims to provide evidence for the gradable nature of projection, and because of that it is also assumed that “at-issue” and “follows from the common ground” are gradient properties as well. I don’t see what warrants that or how it is supported by the data. There are two choices of how gradience can enter the picture, and it seems to me that the first leads to bizarre implications (and is at odds with the two experiments) and that the second does not require gradience at all.

Contrary to what is stated here, our paper does not have the goal of providing evidence for the gradable nature of projection and at-issueness. That evidence, as we now make clearer in sections 1 and 2 of the paper, was provided in prior experimental literature, including in Tonhauser et al 2018. Our paper builds on that literature. Rather, the goal of our paper is to develop an analysis of EAS that accounts for the projection of the prejacent and the interaction between the projectivity of the prejacent and the generalization.

Nevertheless, the findings both of the corpus study and of Experiment 1 further support the finding that projection is not a binary, categorical property of projective content but a gradient one. As shown in Figure 5, there is not just by-condition projection variability but also by-evaluative adjective variability: the prejacent of EAS with some adjectives is more projective than the prejacent of EAS with other adjectives.

On p.8 the paper mentions two ways to interpret gradient projection: as pertaining to the data itself or as quantifying the agent’s uncertainty about the data. The authors go on to say: “we remain agnostic about the underlying interpretation of projectivity as a gradient property, though our discussion of projection variability will be in line with the first interpretation” (p.8). The former choice, which the authors seem to make, leads to things that I find difficult to parse. For example, what would it mean for a prejacent to follow from the common ground to 30%?

The manuscript does not mention percentages, but we can suggest that what this would mean is that the listener believes that the prejacent has a probability of .3 of being true a priori, given their understanding of what is in the common ground.

So, by Non-redundancy, it should be 70% at-issue? And by the Gradient Projection Principle it will project to 30%? Given the complementary nature of prejacent and generalization, does that also mean that the generalization follows from the common ground to 70%, is at-issue to 30%, and projects to 70%? What would be a good example of that?

If the listener a priori assigns the prejacent a .3 probability of being true, it does not follow that the probability of the prejacent being at-issue is .7. This is because the ‘Non-redundancy principle for at-issue content’ allows for a non-linear relationship between the probability with which the prejacent is taken to follow from the common ground and the probability of the Discourse Question being about the prejacent. What the principle states is that if content c_1 is more likely to be taken to follow from the common ground than content c_2 , then the Discourse Question is less likely to be about c_1 than about c_2 . What the function is that maps the difference in prior probability between c_1 and c_2 to the difference in the probability of the Discourse Question being about c_1 vs about c_2 is an interesting empirical question.

Continuing along these lines for the question “And by the Gradient Projection Principle it will project to 30%?”: the Gradient Projection Principle does not stipulate a linear relationship between not-at-issueness and projection, merely a positive one: if content c_1 is more not-at-issue than content c_2 , then the Gradient Projection Principle predicts that content c_1 is more projective than c_2 . And, indeed, Tonhauser et al 2018 found evidence for such a positive relationship, as shown in figures from their paper, one of which we now include in our paper as Figure 2 in section 2.1.

In addition, even if these things can somehow be defined, I don’t see how the two experiments support gradient projection. Both predictors in Experiment 1 are binary, so we cannot set a third intermediate value and observe its effect on the response variable. In Experiment 2, both the predictor and the response are binary, so gradience cannot arise either.

As mentioned above, the goal of the paper is not to provide evidence for the gradient nature of projection or at-issueness. The two experiments were designed, as stated in (20), to provide evidence for two predictions of the analysis:

- (20) The more likely the generalization is taken to be at-issue,
a. ...the more projective the prejacent is. (Experiment1)
b. ...the less likely the prejacent is taken to be at-issue. (Experiment2)

Because the analysis developed in section 2 takes projection and at-issueness to be gradient rather than binary, categorical properties of utterance content, these predictions do, too.

You are right, of course, that Experiments 1 and 2 only involved predictors with two levels: the truth of the generalization is more or less likely to follow from the common

ground. The predictions in (20) thus lead us to expect that the prejacent is more projective (Exp 1) or more at-issue (Exp 2) in one condition than the other. And the experimental findings provide support for that prediction. One could, of course, also run experiments to test the predictions in (20) in which the predictors have more levels.

On the other hand, if we interpret gradience as quantifying uncertainty, we can stick with the binary options for projection and all the supplementary concepts mentioned about. It seems to me that that the paper doesn't offer any evidence for gradient projection after all.

Yes, as stated in the manuscript, on a second interpretation of projectivity as a gradient property of utterance content, speaker commitment is categorical and "projection variability arises from the listener's uncertainty about whether the speaker is committed". As mentioned above, the goal of this manuscript was not to provide evidence for the gradient nature of the projection of utterance content; we build on Tonhauser et al 2018.

#3: As far as I can tell, the examples cited to show that EAS may be read in two different ways (prejacent projects/generalization at-issue, prejacent at-issue/generalization projects) include negation. Since the claim is that EAS have those two entailments by virtue of their lexical semantics, we should be able to see the same kind of "ambiguity" with modal operators as well, for example in EAS like Perhaps/Most likely/Mary thinks Feynman was stupid to dance on the table. This omission should be easy to fix by just citing a few appropriate examples.

Thank you for this suggestion. We have included two examples in footnote 2 to show that non-projection of the prejacent is also observed with other entailment-canceling operators.

#4: There are several unclarities in the semantic analysis in (10), p.8. First, I don't know why VP' takes two arguments in (10a) but only one argument in (10b).

What was examples (10) are now examples (13). In (13b), the VP' takes one argument, the translation of the subject, so that the resulting expression is of type $\langle ev, t \rangle$, i.e., denotes a set of events. The adjective then combines with this event description to express that, at the reference time rt , the degree to which events of this sort are *adj* is higher than the contextual standard of *adj*. If the VP' in (13b) also took an event argument, as it does in (13a), the resulting expression would not denote a set of events anymore, but an expression of type t , which is not an appropriate argument for the adjective.

Second, the prose after (10b) does not corresponds to the simple formula that precedes it. I don't see why we need to talk about degrees, but if these are important for some reason, why not elaborate the formula so it has all the necessary information in it?

Thank you for asking for clarification here. We are following Barker's 2002 analysis of evaluative adjectives here (which included degrees, but was not relativized to times). The

original translation in (13b) did not have degrees in the object language. The modified version of (13b) now included in the manuscript follows Barker's 2002 analysis more closely in including degrees in the object language.

Finally, I don't understand in what sense the generalization is "timeless", given that the generalization of an EAS may get different truth values depending on when it is uttered. Example: The generalization of John was smart to marry Sue (which can be rendered informally as It was smart of John to marry Sue) may be true before the divorce and false after the divorce. This goes back to my second point that the generalization in (10b) needs to be properly spelled out formally, just like the prejacent in (10a) is.

Thank you also for this comment. Barker 2002 did not consider times/temporal reference, but we have now expanded his analysis to do so. Specifically, we have removed the claim that the generalization is timeless and have reformulated (13b) to include the reference time. On this analysis, the generalization of *John was smart to marry Sue* is that, at the past reference time, the degree to which events of John marrying Sue are smart was higher than the contextual standard of smart. This captures the observation that the generalization may be true at one time and false at another.

#5: On p.12 it is stated that Experiment 2 investigates the prediction that "The more likely the generalization is taken to be at-issue, the less likely the prejacent is taken to be at-issue."

In reality, Experiment 2 only manipulates the extent to which the generalization follows from the common ground. These are not the same things, and it takes an extra step to get from common ground status to at-issue status, otherwise we won't need Non-redundancy.

Right, to test the prediction in (20b), that the more likely the generalization is taken to be at-issue, the less likely the prejacent is taken to be at-issue", we need to manipulate the at-issueness of the generalization. According to the 'Non-redundancy principle for at-issue content', the more the truth of content follows from the common ground, the less likely it is to be at-issue. Thus, assuming this principle, Experiment 2 manipulates the at-issueness of the generalization by manipulating the extent to which the truth of the content follows from the common ground. We now clarify in the text before (26) that we manipulated the at-issueness of the generalization by manipulating the extent to which it is taken to follow from the common ground, and how this manipulation relies on the 'Non-redundancy principle for at-issue content'.

#6: "the prejacent and the generalization of EASs are independent of one another in that neither is a precondition for the truth of the other. Consider (1) Feynman was stupid to dance on the table: it is neither necessary for Feynman to have danced on the table in order for the generalization of this EAS to be true, nor is it necessary for the generalization to be true in order for Feynman to have danced on the table."

The latter statement is clearly true, but I was surprised to read the former statement. If generalization and prejacent are indeed logically independent, as claimed, what would be

a context in which the generalization of a (unmodified past tense) EAS is true but the prejacent is false?

Here is a context in which the prejacent of (1) is false but the generalization of (1) is true: at the past reference time, Feynman did not dance on the table (prejacent is false) but the degree to which events of Feynman dancing on the table are stupid was higher at that reference time than the contextual standard of stupid (generalization is true), i.e., if Feynman had danced on the table, that would have been stupid. We have clarified the relevant passage in section 5.

Review B:

In my judgment, the main proposal of the paper, i.e., the question-based analysis of the projective pattern in “evaluative adjective sentences”, is incomplete.

The key idea, according to the authors, is that “projection is not derived from a lexical specification but rather derived from projective content being backgrounded or not-at-issue”.

For one thing, it is not clear whether this is meant to be a generic statement about projective content or rather about specific cases, such as “the projection of the prejacent of the EAS”. I think, and the authors’ reply and new manuscript suggest that they agree, that the first interpretation is rather hard to maintain.

Yes, the statement is meant to only apply to the prejacent and the generalization of EAS. We have revised relevant passages in the abstract and the introduction to make this very clear. For instance, the abstract now states that “We argue against such analyses **of the prejacent...**” (bold-faced material added). That we do not argue that a question-based analysis is appropriate for all classes of projective content is now also mentioned in sections 1 and 5, as well as in footnote 6.

A key observation here is that in (1), either entailment #1 or #2 may project, and exactly one of them must project.

(1) John was smart to buy a Prius.

entailment #1 (prejacent-proposition): John bought a Prius

entailment #2 (generalization): John’s buying a Prius would be smart.

Typically it is entailment #1 that projects, but (for some speakers at least) when the DQ is “Did John buy a Prius?”, it is entailment #2 that (is backgrounded and therefore) projects. But the same does not happen in most other cases. Unlike (1), (2) cannot be a felicitous answer to “Did John buy a Prius?” – it cannot be interpreted in such a way that the prejacent is non-projective. In other words, the CI-status of the prejacent-proposition is hardwired to the construction.

(2) It was smart of John to buy a Prius.

entailment #1 (= **CI**): John bought a Prius

entailment #2 (= **proffered content**): John's buying a Prius would be smart.

Right: the prejacent of the 'It was Adj of NP to VP' construction seems to differ in its properties from that of the evaluative adjective sentences we investigate in this manuscript, which are of the form 'NP was Adj to VP'. As we noted in the previous response letter, we do not assume that the projectivity of all content is derived from its discourse properties. The current version of the manuscript expands on the relevant footnote (now footnote 6), to state explicitly that the prejacent of the 'It was Adj of NP to VP' construction may be best analyzed as a lexically specified presupposition.

Similarly, in (3B), it is entailment #1 that is "foregrounded", providing information that helps solving the DQ. However it does not make entailment #2 ('John's car is a Prius') projective; evidence is that (3B') does not entail that John's car is a Prius.

(3) A: What color is John's car? B: It is a blue Prius.

entailment #1: John's car is blue.

entailment #2: John's car is a Prius.

B': It may be a blue Prius.

Right, B' does not entail that John's car is a Prius. But the important question is whether the content that John's car is a Prius projects from B', i.e., whether the speaker of B' may be taken to be committed to John's car being a Prius. As far as we know, the projection of such content has not yet been investigated experimentally. According to our intuitions, at least, it is possible for the proposition that John's car is a Prius to project from an utterance of B', namely when *blue* carries the nuclear pitch accent. This prosody is not, of course, felicitous in response A's question, but this illustration at the very least shows that the proposition that John's car is Prius can project.

So, the variable projective pattern of sentence (1) must have a root in the lexically (constructionally), i.e. conventionally, specified semantics of the EAS.

Unfortunately, we do not understand your argument here: how does it follow from the observations regarding the related but different construction in (2) and the completely different example in (3) that the prejacent of EAS "must have a root in the lexically...specified semantics of the EAS"? As discussed in the paper, analyses that assume that the prejacent is lexically specified to project do not make empirically adequate predictions.

If the semantics of (2) is to be phrased as in (4), that of (1) will look like (5) ("variable entailment" is a tentative label applied to a semantic component that can end up either being proffered or conventionally implicated).

4. (4) It was smart of John to buy a Prius.

CI: ‘John bought a Prius’

Proffered Content: ‘John’s buying a Prius would be smart’

5. (5) John was smart to buy a Prius.

Variable Entailment #1: ‘John bought a Prius’

Variable Entailment #2: ‘John’s buying a Prius would be smart’

(additional conditions) Exactly one of VE #1 or VE #2 is proffered and the other is conventionally implicated. There is a general (rather strong), conventionalized preference for the interpretation where VE #1 is a CI. The DQ constrains which interpretation the sentence receives.

The analysis of EAS you sketch here is similar to that of Karttunen et al 2014, except that their analysis does not rely on the DQ. In our manuscript, we engage with Karttunen et al’s analysis in sections 1 and 5, and it is possible that some of the arguments against Karttunen et al’s analysis would also apply to a fully fleshed out analysis along the lines of what you sketch here.

The authors claim that it does not need to be lexically specified which semantic component of the AES is conventionally implicated (projective), and its variable projective pattern follows from the principle such that the material that is not directly relevant to the DQ projects. But for this to work, *we need to lexically or constructionally specify, beforehand, that the two components of the AES are variable entailments, rather than hardwired proffered content or CI.* (In (5), the DQ may play an important role but it only constitutes part of the additional conditions.) This is an *idiosyncratic* (and peculiar/intriguing) feature of the AES construction, absent in most if not all other constructions including the “It is *Adj.* of *N* to *VP*” construction, and which needs to be specified at the level of conventional meaning.

You write that “we need to lexically or constructionally specify, beforehand, that the two components of the EAS are variable entailments, rather than hardwired proffered content or CI”. Our manuscript states that the prejacent and the generalization are lexical entailments, and it also states that both can be at-issue and not at-issue. We also explicitly say that neither of them is a lexically specified presupposition. So, as far as we can tell, our manuscript does exactly what you say we need to do.

And, yes, EASs are different from other constructions with projective content, as we make clear in sections 1 and 5, and in footnote 6.

As such, it is not clear whether the “question-based account” is much less costly than a simple polysemy account along the lines of (6).

- (6) John was smart to buy a Prius.

Sense #1 (preferred/more common one)

CI: ‘John bought a Prius’

Proffered Content: ‘John’s buying a Prius would be smart’

Sense #2 (marked/marginal one)

CI: ‘John’s buying a Prius would be smart’

Proffered Content: ‘John bought a Prius’

As noted in my previous review comments (with the *mole* example), it is independently acknowledged fact that DQ’s, among other contextual factors, contribute much to sense disambiguation. In a way, (6) is more parsimonious than (5) in not involving direct reference to the DQ. (For that matter, it is questionable if it is desirable to include it in (5); the last sentence of (5) may as well be left out, assuming that the relevant polysemy is resolved in the same way as polysemy in general is.)

The authors should provide a semantic formulation of the EAS that is at least as detailed and systematic as (5), and explain what exactly we gain by adopting it instead of an alternative like (6). (I see that the question-based approach to projection has gained some positive reaction and popularity, but I for one remain skeptic about it; the criticism I raised above probably will apply to the question-based approach in general.) The notion of Variable Entailment (or whatever one calls it) will be an interesting idea to pursue, while it is yet to be seen whether we should admit the existence of such a thing. I am highly doubtful that reference to the DQ is an essential ingredient in the proper account of the observed phenomena.

[Thank you for this comment. As mentioned at the beginning of this letter, we have added in section 2.4 a comparison of our analysis to the ambiguity analysis presented in Karttunen et al 2014.](#)

Review C:

The current version of the article represents a significant improvement from the previous version. All the issues that I pointed out then have been addressed to my satisfaction, and I therefore recommend that it be accepted for publication. In particular, I appreciated the degree to which the authors have streamlined their prose and polished the overall organization of the paper, which is now considerably easier to follow.

[Thank you! We are glad to hear this.](#)

Nevertheless, there are two issues that I think deserve some discussion. None of them is necessary for publication, but it’d be interesting if the authors could perhaps provide some broad discussion of these in the final part of the paper, perhaps replacing the discussion of inter-speaker variation, which I didn’t find particularly informative. First, what is the role of evaluativity *per se* in determining the observed pattern — especially as it relates to gradability more broadly? As far as I understand, the authors discuss the semantics of the evaluativity only in terms of a relation between two degrees — the one instantiated by an individual and at the one associated with a contextual standard. While I

assume this approach is feasible to capture the semantics of evaluativity, note that a very similar semantics applies to gradable adjectives more broadly. What about *fast* for example?

(1) Mary wasn't fast to grab the glass before it fell.

Is the sentence above grammatical in the first place? If so, what is its projectivity profile? Is it as discourse-sensitive as the one of *smart* or *stupid*? Note that *fast* might still be construed as evaluative, though it doesn't seem to be of the same kin as the adjectives tested in the experiment—at the very least, it encodes a mono-dimensional core, contrary to the multi-dimensional nature of the adjectives discussed in the paper. In any case, it would be nice to provide some discussion—and minimally to describe how the authors went about determining the evaluative status of the attributes tested in the study.

Yes, it seems that sentences like your (1) can be judged to be acceptable. See, for instance, the naturally occurring example in (1'), which seems to mean that Stewart got up but didn't do so in a fast manner.

(1') [Malcom Stewart] just wasn't fast to get up after crashing.
<https://motocrossactionmag.com/10-things-you-need-to-know-about-anaheim-i/>

We do not know whether such examples can also have a reading according to which the prejacent does not project: native speakers we have consulted cannot get such interpretations and we have not encountered such examples in naturally occurring data, but we also haven't conducted an experimental investigation.

The adjective *fast* is not an evaluative adjective because it does not convey the speaker's attitude/evaluation on the denotation of the subject noun phrase, i.e., an utterance of a sentence like *Mary was fast* does not convey the speaker's attitude on Mary (it may have been good that she was fast or not). We have added a brief characterization of evaluative adjectives in footnote 1.

Second, what implications do these results have for the study of presupposition projection more generally? It seems like the approach advocated by the author—which crucially suggests that projection is better captured as a discourse-level phenomenon than a lexical one—could have serious implications for the account of other instances of projectivity above and beyond the domain of evaluative adjectives. This is obviously outside the scope of the paper, and thus I'm not expecting the authors to address the issue in an exhaustive manner. Nevertheless, it might be interesting to at least make some suggestive comments as to whether the discourse-sensitivity of projectivity is limited to the phenomenon considered in the paper, or extends more broadly.

Thank you for this suggestion. We have included a paragraph in section 5 that discusses the implications for investigations of projective content.